

# **ESSENTIALS OF MACHINE LEARNING (EML)**

## **PBL- ACTIVITY**

**TITLE : HEART DISEASE**

**TEAM:**

**23EG107F30- YASANI SHREYAS REDDY**

**23EG107F35- PRANATHI KATAPALLY**

**23EG107F39 - KATTA AKASH REDDY**

**23EG107F52 - G.SRI SAHASRA**

**23EG107F53 - G.SRI SAHAJA**

# DATA SET

Data set link : #!/bin/bash  
kaggle datasets download  
shrutikubade/heart-diseases-dataset

- ❖ Number of records : 303 records
- ❖ Number of columns :14
- ❖ Number of classes: 2

Class(0): No heart disease

Class (1): Heart disease present

**The Dependent features or the target variable includes:** No heart disease and heart disease presence i.e 0 and 1. Dependent variable is ‘target’.

**The independent features includes:** ‘age’, ‘sex’, ‘cp’, ‘trestbps’, ‘chol’, ‘fbs’, ‘restecg’, ‘thalach’, ‘exang’, ‘oldpeak’, ‘slope’, ‘ca’, ‘thal’.



# ATTRIBUTE INFORMATION

**1. age** : Age of the patient (in years)

**2. sex**: Sex of the patient (1 = male, 0 = female)

**3. cp**: Chest pain type (categorical: 0–3)

**4. trestbps**: Resting blood pressure

**5. chol**: Serum cholesterol

**6. fbs** – Fasting blood sugar >120 mg/dl (1 = true, 0 = false)

**7. restecg**: Resting electrocardiographic results

- 0: Normal
- 1: Having ST-T wave abnormality
- 2: Showing probable/definite left ventricular hypertrophy

**8. thalach:** Maximum heart rate achieved

**9. exang:** Exercise-induced angina (1 = yes, 0 = no)

**10. oldpeak:** ST depression induced by exercise relative to rest (float value)

**11. slope:** The slope of the peak exercise ST segment

- 0: Upsloping
- 1: Flat
- 2: Downsloping

**12.ca** – Number of major vessels (0–3) colored by fluoroscopy



### **13.thal – Thalassemia status**

- 1: Normal
- 2: Fixed defect
- 3: Reversible defect

### **14. target (Dependent Variable) – Presence of heart disease**

- 0: No heart disease
- 1: Heart disease present



# IMPORTING THE PACKAGES

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

# IMPORTING THE DATASET

```
df=pd.read_csv("/content/Day9_Heart_Disease_Data.csv")
df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

# MISSING VALUES INFORMATION

- ★ There are **no missing values** in any of the 14 columns.
- ★ All 303 records are **complete**.
- ★ No imputation
- ★ Dataset is already clean and ready for preprocessing (encoding, scaling, etc.).



```
df.isnull().sum()
```

	0
age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0
oldpeak	0
slope	0
ca	0
thal	0
target	0

dtype: int64

# UNDERSTANDING THE DATASET

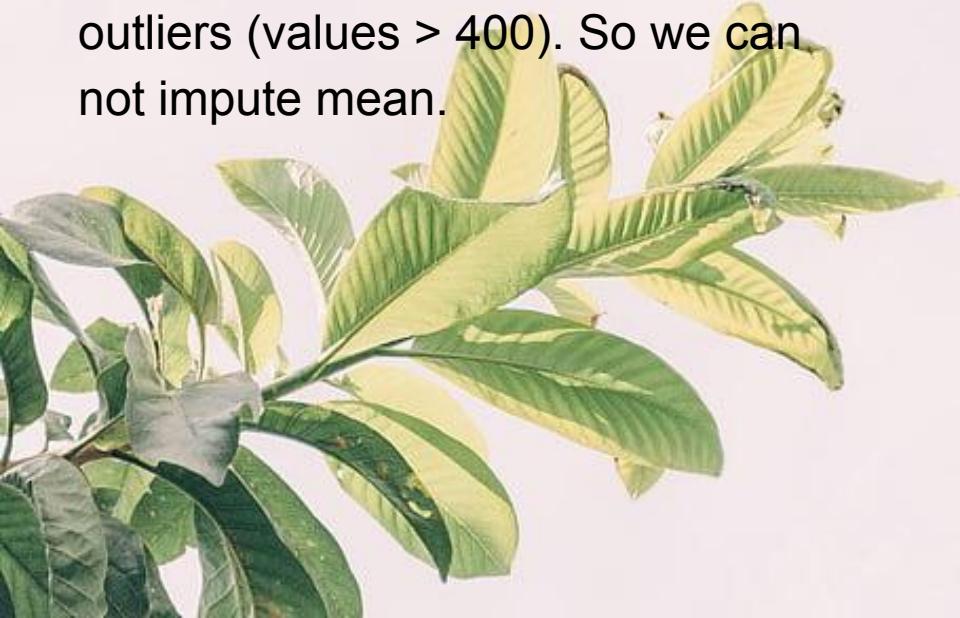
```
columns=['cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal','age', 'trestbps', 'chol', 'thalach', 'oldpeak']
unique_values={col:df[col].unique() for col in columns}
unique values
```

The features sex, cp, fbs, restecg, exang, slope, ca, thal are Categorical Feature and age, trestbps, chol, thalach, oldpeak are Numerical Features.

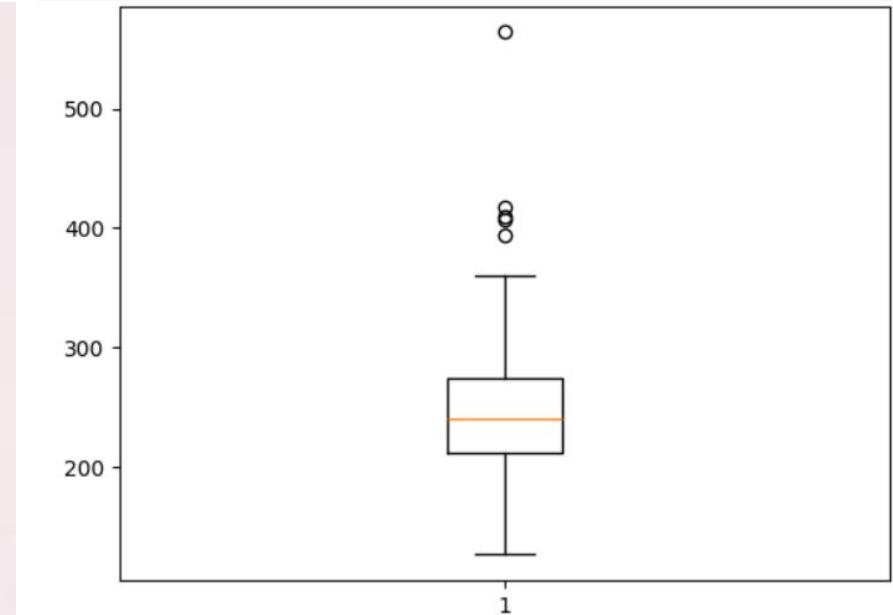
```
{'cp': array([3, 2, 1, 0]),
 'fbs': array([1, 0]),
 'restecg': array([0, 1, 2]),
 'exang': array([0, 1]),
 'slope': array([0, 2, 1]),
 'ca': array([0, 2, 1, 3, 4]),
 'thal': array([1, 2, 3, 0]),
 'age': array([63, 37, 41, 56, 57, 44, 52, 54, 48, 49, 64, 58, 50, 66, 43, 69, 59,
 42, 61, 40, 71, 51, 65, 53, 46, 45, 39, 47, 62, 34, 35, 29, 55, 60,
 67, 68, 74, 76, 70, 38, 77]),
 'trestbps': array([145, 130, 120, 140, 172, 150, 110, 135, 160, 105, 125, 142, 155,
 104, 138, 128, 108, 134, 122, 115, 118, 100, 124, 94, 112, 102,
 152, 101, 132, 148, 178, 129, 180, 136, 126, 106, 156, 170, 146,
 117, 200, 165, 174, 192, 144, 123, 154, 114, 164]),
 'chol': array([233, 250, 204, 236, 354, 192, 294, 263, 199, 168, 239, 275, 266,
 211, 283, 219, 340, 226, 247, 234, 243, 302, 212, 175, 417, 197,
 198, 177, 273, 213, 304, 232, 269, 360, 308, 245, 208, 264, 321,
 325, 235, 257, 216, 256, 231, 141, 252, 201, 222, 260, 182, 303,
 265, 309, 186, 203, 183, 220, 209, 258, 227, 261, 221, 205, 240,
 318, 298, 564, 277, 214, 248, 255, 207, 223, 288, 160, 394, 315,
 246, 244, 270, 195, 196, 254, 126, 313, 262, 215, 193, 271, 268,
 267, 210, 295, 306, 178, 242, 180, 228, 149, 278, 253, 342, 157,
 286, 229, 284, 224, 206, 167, 230, 335, 276, 353, 225, 330, 290,
 172, 305, 188, 282, 185, 326, 274, 164, 307, 249, 341, 407, 217,
 174, 281, 289, 322, 299, 300, 293, 184, 409, 259, 200, 327, 237,
 218, 319, 166, 311, 169, 187, 176, 241, 131]),
 'thalach': array([150, 187, 172, 178, 163, 148, 153, 173, 162, 174, 160, 139, 171,
 144, 158, 114, 151, 161, 179, 137, 157, 123, 152, 168, 140, 188,
 125, 170, 165, 142, 180, 143, 182, 156, 115, 149, 146, 175, 186,
 185, 159, 130, 190, 132, 147, 154, 202, 166, 164, 184, 122, 169,
 138, 111, 145, 194, 131, 133, 155, 167, 192, 121, 96, 126, 105,
 181, 116, 108, 129, 120, 112, 128, 109, 113, 99, 177, 141, 136,
 97, 127, 103, 124, 88, 195, 106, 95, 117, 71, 118, 134, 90]),
 'oldpeak': array([2.3, 3.5, 1.4, 0.8, 0.6, 0.4, 1.3, 0., 0.5, 1.6, 1.2, 0.2, 1.8,
 1., 2.6, 1.5, 3., 2.4, 0.1, 1.9, 4.2, 1.1, 2., 0.7, 0.3, 0.9,
 3.6, 3.1, 3.2, 2.5, 2.2, 2.8, 3.4, 6.2, 4., 5.6, 2.9, 2.1, 3.8,
 4.4])}
```

# UNDERSTANDING THE DATASET

There are many outliers in Cholesterol (chol) has the most outliers (values > 400). So we can not impute mean.



[25] `import matplotlib.pyplot as plt  
plt.boxplot(df['chol'])`



# INFORMATION ABOUT THE DATASET

Now we need to convert all the categorical variables into numerical using encoding.



ds df.info()

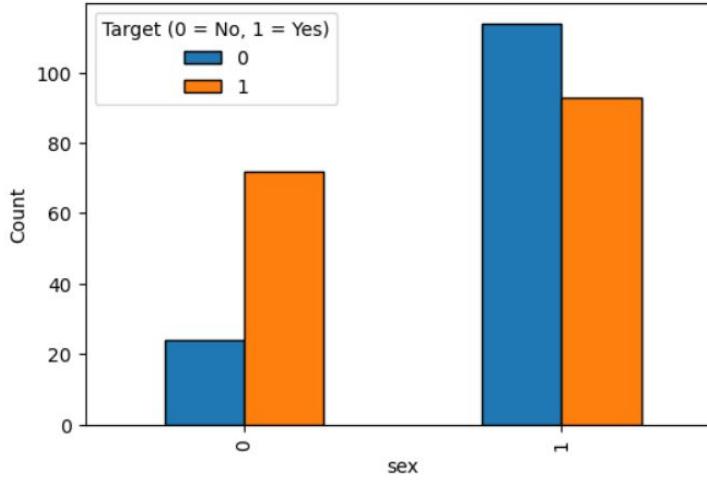
```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         303 non-null    int64  
 1   sex         303 non-null    int64  
 2   cp          303 non-null    int64  
 3   trestbps    303 non-null    int64  
 4   chol        303 non-null    int64  
 5   fbs         303 non-null    int64  
 6   restecg     303 non-null    int64  
 7   thalach     303 non-null    int64  
 8   exang       303 non-null    int64  
 9   oldpeak     303 non-null    float64 
 10  slope       303 non-null    int64  
 11  ca          303 non-null    int64  
 12  thal        303 non-null    int64  
 13  target      303 non-null    int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

# BAR PLOTS (Categorical variables vs target variables)

```
▶ import pandas as pd
import matplotlib.pyplot as plt
cat_cols = ["sex", "cp", "fbs", "restecg", "exang", "slope", "ca", "thal"]
for col in cat_cols:
    plt.figure(figsize=(6,4))
    pd.crosstab(df[col], df["target"]).plot(kind="bar", figsize=(6,4), edgecolor="black")
    plt.title(f"Bar Plot of {col} vs Target")
    plt.xlabel(col)
    plt.ylabel("Count")
    plt.legend(title="Target (0 = No, 1 = Yes)")
    plt.show()
```

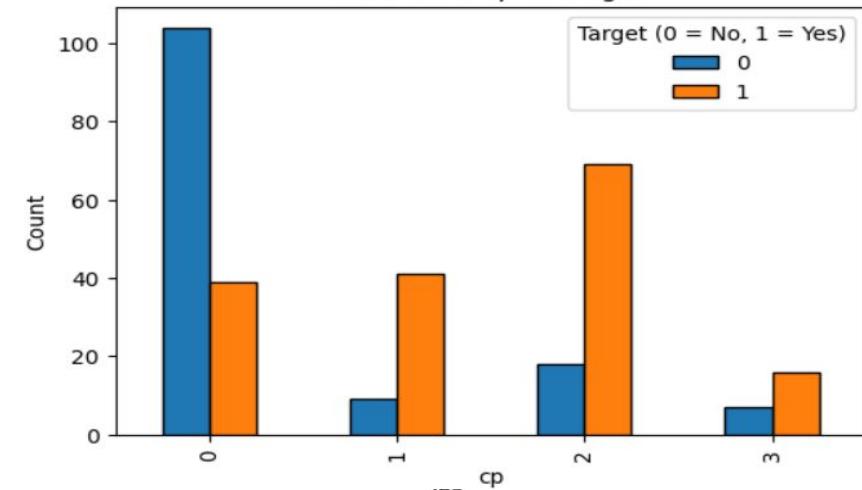
⤵ <Figure size 600x400 with 0 Axes>

Bar Plot of sex vs Target



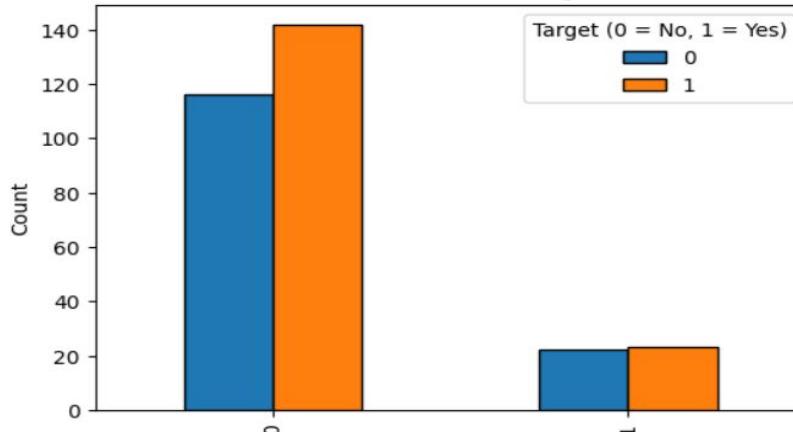
⤵ <Figure size 600x400 with 0 Axes>

Bar Plot of cp vs Target



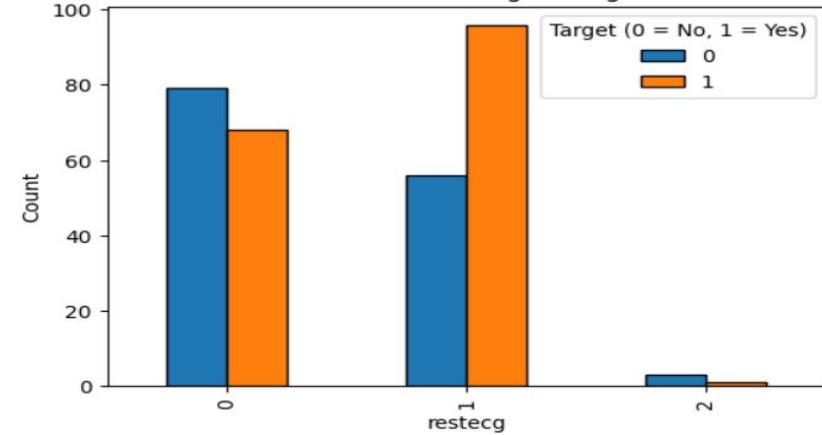
⤵ <Figure size 600x400 with 0 Axes>

Bar Plot of fbs vs Target



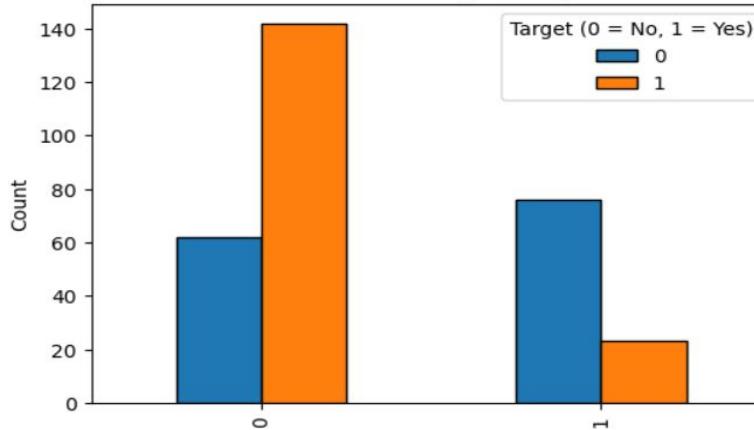
⤵ <Figure size 600x400 with 0 Axes>

Bar Plot of restecg vs Target



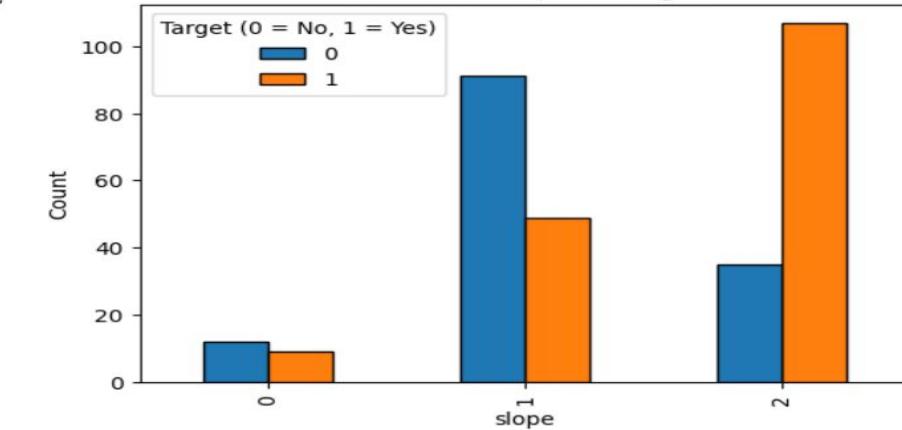
<Figure size 600x400 with 0 Axes>

### Bar Plot of exang vs Target



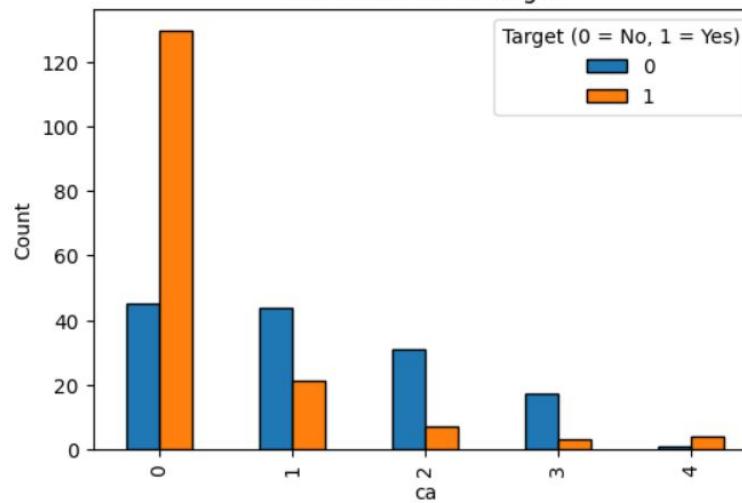
<Figure size 600x400 with 0 Axes>

### Bar Plot of slope vs Target



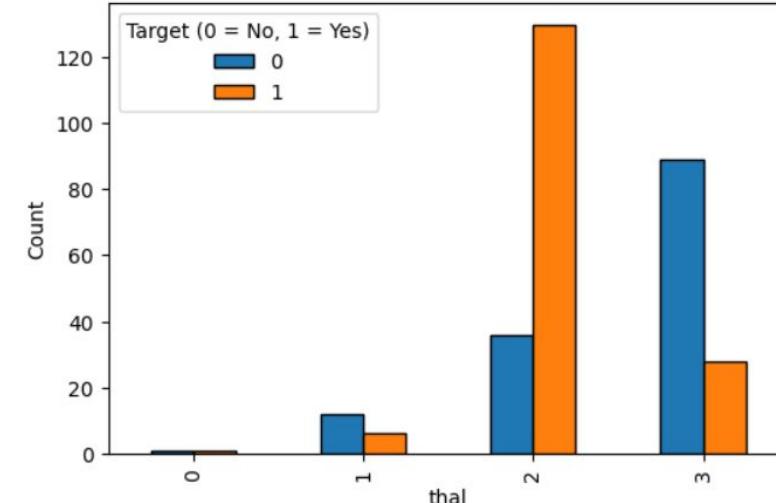
<Figure size 600x400 with 0 Axes>

### Bar Plot of ca vs Target



<Figure size 600x400 with 0 Axes>

### Bar Plot of thal vs Target



# EXPLORATORY DATA ANALYSIS



df.describe()

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000

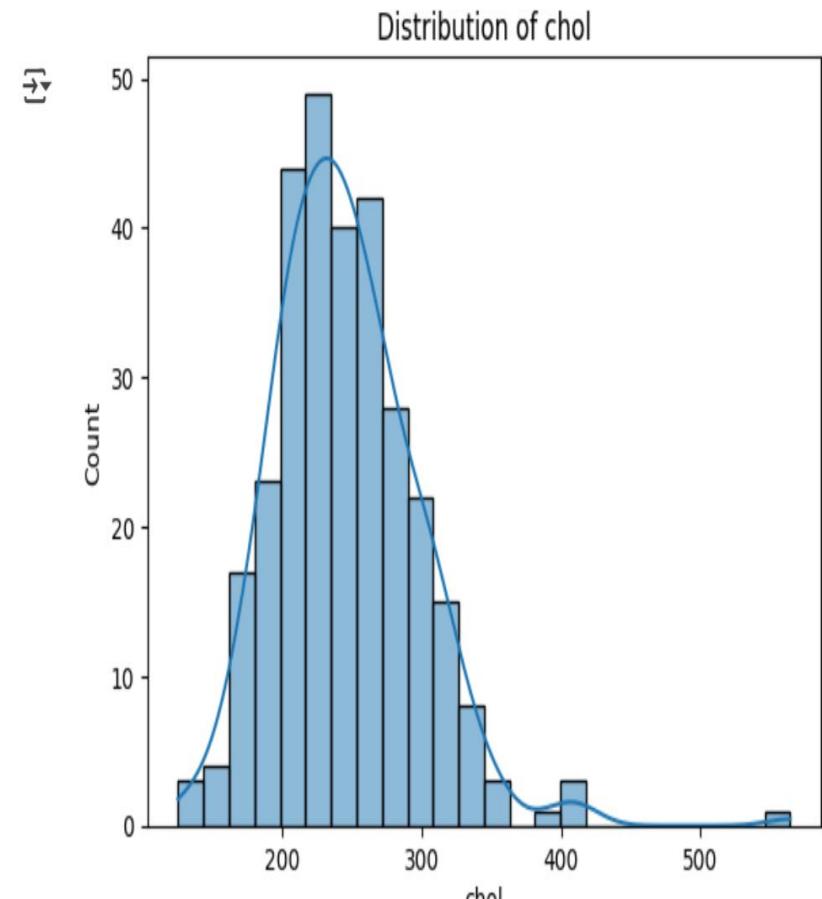
oldpeak	slope	ca	thal	target
3.000000	303.000000	303.000000	303.000000	303.000000
1.039604	1.399340	0.729373	2.313531	0.544554
1.161075	0.616226	1.022606	0.612277	0.498835
0.000000	0.000000	0.000000	0.000000	0.000000
0.000000	1.000000	0.000000	2.000000	0.000000
0.800000	1.000000	0.000000	2.000000	1.000000
1.600000	2.000000	1.000000	3.000000	1.000000
6.200000	2.000000	4.000000	3.000000	1.000000



This data provides summary of mean, count, min, standard deviation.

# UNIVARIATE ANALYSIS

```
import seaborn as sns  
sns.histplot(df['chol'],kde=True)  
plt.title('Distribution of chol')  
plt.show()
```



# BI-VARIATE ANALYSIS

Understanding the correlation between different features.



Os

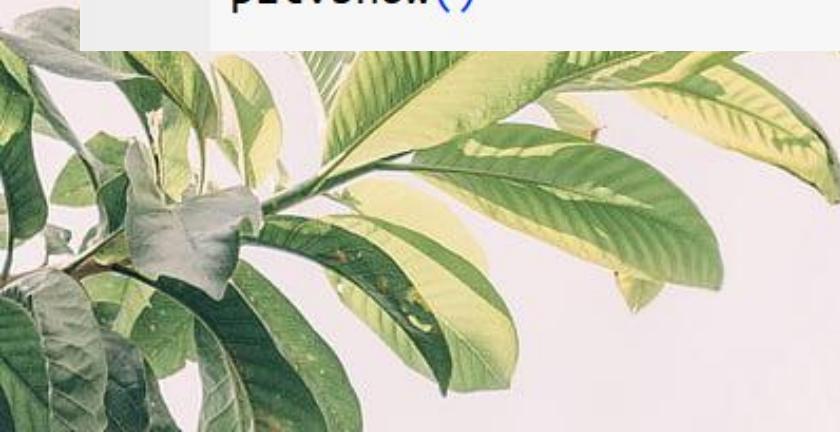
[56]

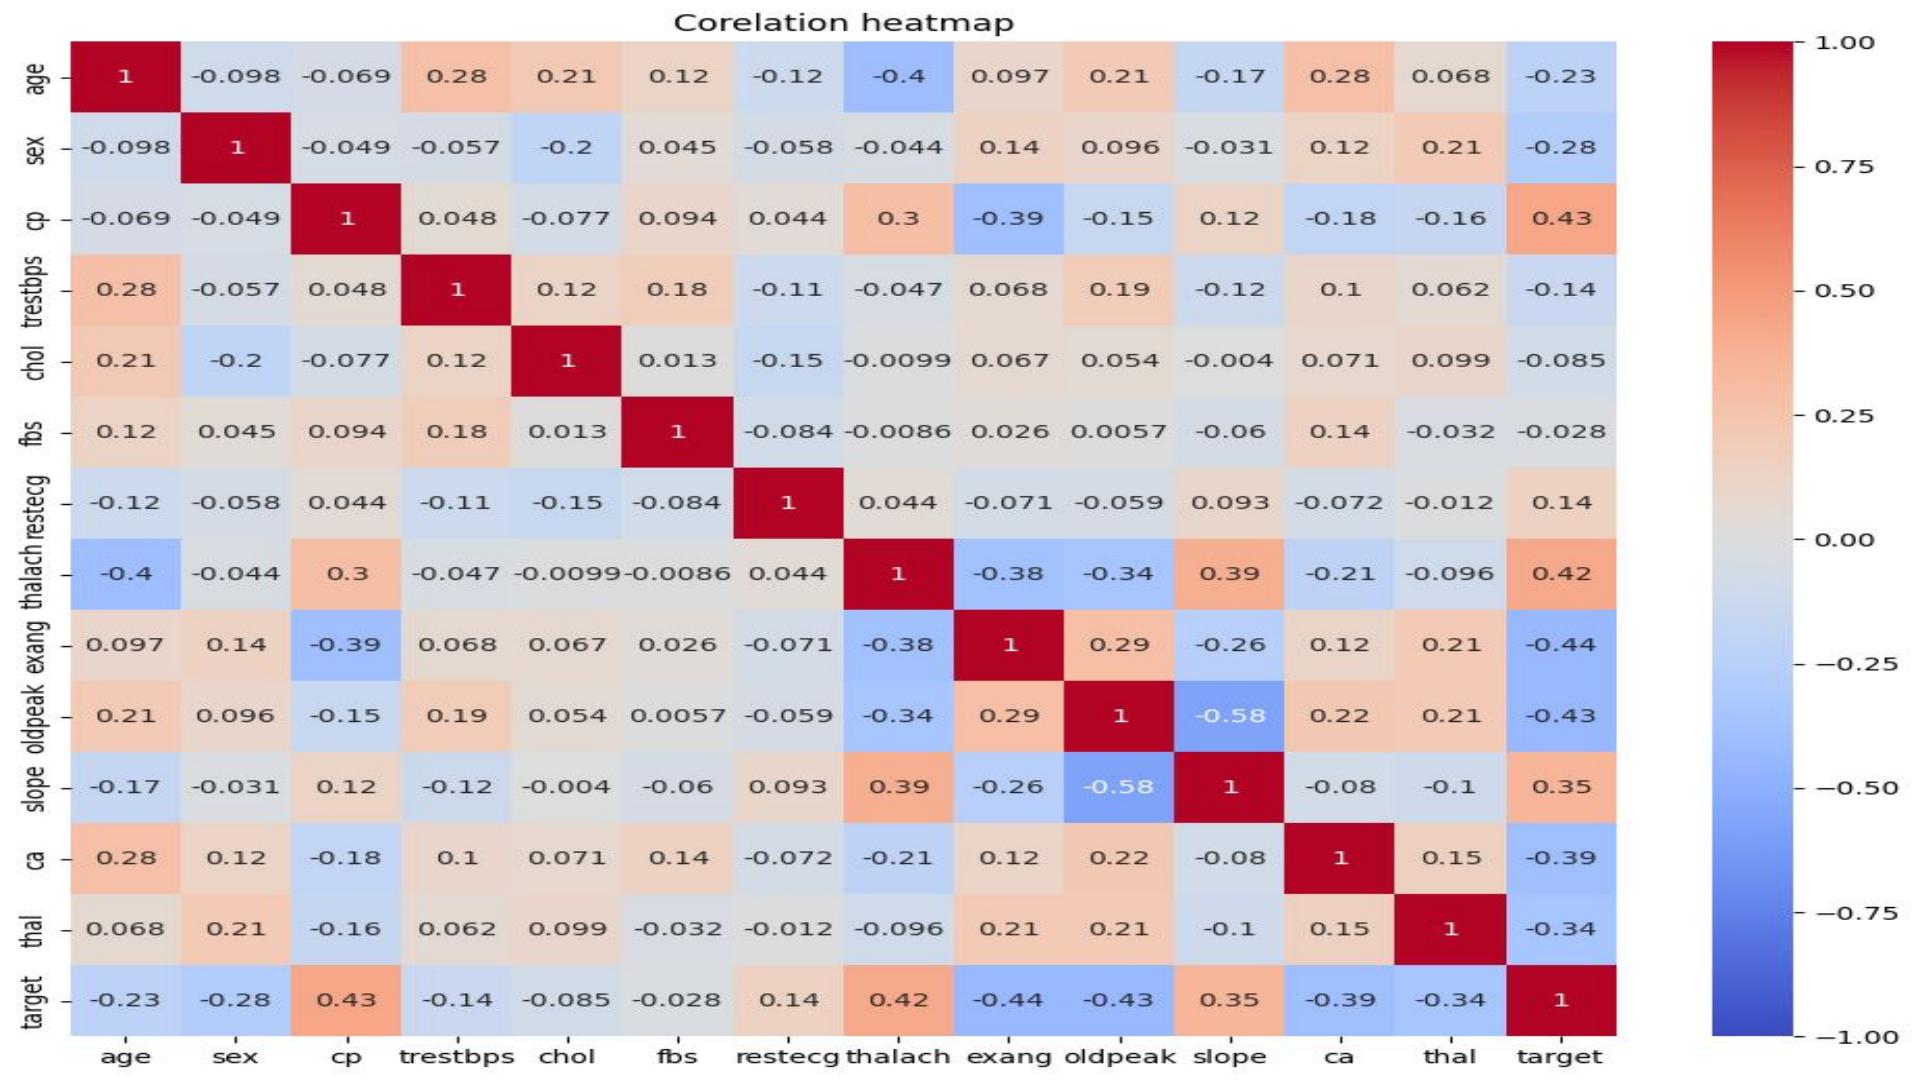
df.corr()

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522	0.096801	0.210013	-0.168814
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020	0.141664	0.096093	-0.030711
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762	-0.394280	-0.149230	0.119717
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698	0.067616	0.193216	-0.121475
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940	0.067023	0.053952	-0.004038
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567	0.025665	0.005747	-0.059894
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123	-0.070733	-0.058770	0.093045
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000	-0.378812	-0.344187	0.386784
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378812	1.000000	0.288223	-0.257748
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.344187	0.288223	1.000000	-0.577537
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.386784	-0.257748	-0.577537	1.000000

# CORRELATION HEATMAP

```
▶ plt.figure(figsize=(12,10))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Corelation heatmap')
plt.show()
```





# PAIR PLOT

```
[61] sns.pairplot(df,hue='chol')
```



# DIVIDING THE DATASET INTO DEPENDENT AND INDEPENDENT FEATURES

```
[66] import pandas as pd
     # Dependent feature (Target)
     y = df.loc[:, 'target']    # target column only

     # Independent features (all columns except target)
     X = df.iloc[:, df.columns != 'target']

     print("Dependent Feature (y):")
     print(y.head())

     print("\n Independent Features (X):")
     print(X.head())
```

## DEPENDENT FEATURES

Dependent Feature (y):

```
0    1  
1    1  
2    1  
3    1  
4    1
```

Name: target, dtype: int64

## INDEPENDENT FEATURES

Independent Features (x):

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

ca thal

0	0	1
1	0	2
2	0	2
3	0	2
4	0	2

# SPLITTING THE DATASET INTO TRAINING AND TESTING DATA

```
✓ Ds ➔ from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)  
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

→ ((242, 13), (61, 13), (242,), (61,))

# STANDARDIZATION

```
✓ [82] from sklearn.preprocessing import StandardScaler  
      scaler=StandardScaler()  
      X_train=scaler.fit_transform(X_train)  
      X_train
```

```
→ array([[-1.35679832,  0.72250438,  0.00809909, ...,  0.95390513,  
        -0.68970073, -0.50904773],  
        [ 0.38508599,  0.72250438, -0.97189094, ...,  0.95390513,  
        -0.68970073,  1.17848036],  
        [-0.92132724,  0.72250438,  0.98808912, ..., -0.69498803,  
        -0.68970073, -0.50904773],  
        ...,  
        [ 1.58263146,  0.72250438,  1.96807914, ..., -0.69498803,  
        0.32186034, -0.50904773],  
        [-0.92132724,  0.72250438, -0.97189094, ...,  0.95390513,  
        -0.68970073,  1.17848036],  
        [ 0.92942484, -1.38407465,  0.00809909, ...,  0.95390513,  
        1.33342142, -0.50904773]])
```

```
X_test=scaler.transform(X_test)  
X_test
```

```
-1.04610909e+00, -2.29340266e-01, 1.47790748e+00,  
-1.93787048e-01, -6.94988026e-01, 3.21860343e-01,  
1.17848036e+00],  
[-7.03591701e-01, 7.22504380e-01, -9.71890936e-01,  
-2.14066346e-02, 1.73802865e-01, 2.60891771e+00,  
-1.04610909e+00, -5.18701733e-03, 1.47790748e+00,  
-9.20864033e-01, 9.53905134e-01, 1.33342142e+00,  
1.17848036e+00],  
[-1.03019501e+00, 7.22504380e-01, -9.71890936e-01,  
6.93132066e-01, 1.17975754e+00, -3.83300706e-01,  
-1.04610909e+00, -1.39678967e-01, 1.47790748e+00,  
-9.20864033e-01, -6.94988026e-01, 2.34498250e+00,  
1.17848036e+00],  
[-1.46566609e+00, 7.22504380e-01, 8.09909113e-03,  
2.76317824e-01, -8.32151805e-01, -3.83300706e-01,  
8.43132697e-01, -8.12138713e-01, -6.76632341e-01,  
-9.20864033e-01, -6.94988026e-01, -6.89700735e-01,  
-2.19657581e+00],  
[ 6.02821534e-01, -1.38407465e+00, 9.88089118e-01,  
-1.68866360e+00, 1.35058003e+00, -3.83300706e-01,  
8.43132697e-01, 4.43119480e-01, -6.76632341e-01,  
-9.20864033e-01, 9.53905134e-01, 3.21860343e-01,  
-5.09047728e-01],  
[ 3.85085995e-01, 7.22504380e-01, -9.71890936e-01,  
-3.19131093e-01, 1.00893504e+00, -3.83300706e-01,  
-1.04610909e+00, 9.36256628e-01, -6.76632341e-01,  
-9.20864033e-01, 9.53905134e-01, 1.33342142e+00,  
1.17848036e+00]
```

# FITTING THE MODEL

```
[87] from sklearn.linear_model import LinearRegression  
regression=LinearRegression()  
regression.fit(X_train,y_train)
```



LinearRegression



LinearRegression()

# PREDICTING THE OUTPUT



```
y_pred=regression.predict(X_test)  
y_pred
```

```
array([ 0.20461609,  0.62932079,  0.71803669,  0.04394647,  0.93859572,  
       0.82919673,  0.55248505, -0.31123783, -0.12686594,  0.50190798,  
       0.67329639,  0.22858348,  0.83576074,  0.10160178,  1.1070706 ,  
       0.90376346,  1.08827421,  0.21347573, -0.14247183, -0.02892816,  
       0.60715419, -0.05665451,  0.35685167,  0.68434807,  0.88713565,  
       0.60700687,  0.81935794,  0.53123747, -0.10013016,  0.8970951 ,  
       0.05884203,  0.07332101, -0.14826886,  0.24195448,  0.69307817,  
       0.19714348,  0.66820272,  0.77828265,  0.68233201,  0.75469071,  
       0.49240619,  0.64012072,  0.75759582,  0.67302635,  0.74556992,  
      -0.1916457 ,  0.66356747,  0.91773248,  0.22766617, -0.00575935,  
      0.15521194, -0.13222102,  0.78206754,  1.04130301,  0.33798087,  
      -0.2796243 ,  0.11850097,  0.94519132,  0.01120021, -0.218084 ,  
      0.15742582])
```

# CONFUSION MATRIX



```
from sklearn.metrics import confusion_matrix  
cm_logistic = confusion_matrix(y_test, y_pred_logistic)  
print("Confusion Matrix:")  
display(cm_logistic)
```



Confusion Matrix:  
array([[25, 4],  
 [ 5, 27]])

# CLASSIFICATION REPORT

```
✓ [102] from sklearn.metrics import classification_report  
      print(classification_report(y_test, y_pred_logistic))
```

	precision	recall	f1-score	support
0	0.83	0.86	0.85	29
1	0.87	0.84	0.86	32
accuracy			0.85	61
macro avg	0.85	0.85	0.85	61
weighted avg	0.85	0.85	0.85	61



Thank  
you